

INTRODUCTION TO PYTORCH, A DEEP LEARNING LIBRARY

**Exercise 1****Getting started with PyTorch tensors**

Tensors are PyTorch's core data structure and the foundation of deep learning. They're similar to NumPy arrays but have unique features.

Here you have a Python list named `temperatures` containing daily readings from two weather stations. Try converting this into a tensor!

Instructions

- Begin by importing `torch`.
- Create a tensor from the Python list `temperatures`.

```
# Import PyTorch
import torch

temperatures = [[72, 75, 78], [70, 73, 76]]

# Create a tensor from temperatures
temp_tensor = torch.tensor(temperatures)

print(temp_tensor)

tensor([[72, 75, 78],
        [70, 73, 76]])
```

Exercise 2**Checking and adding tensors**

While collecting temperature data, you notice the readings are off by two degrees. Add two degrees to the `temperatures` tensor after verifying its shape and data type with `torch` to ensure compatibility with the adjustment tensor.

The `torch` library and the `temperatures` tensor are loaded for you.

Instructions 1

- Display the shape of the adjustment tensor.
- Display the data type of the adjustment tensor.

```
adjustment = torch.tensor([[2, 2, 2], [2, 2, 2]])

# Display the shape of the adjustment tensor
print("Adjustment shape:", adjustment.shape)

# Display the type of the adjustment tensor
print("Adjustment type:", adjustment.dtype)

print("Temperatures shape:", temp_tensor.shape)
print("Temperatures type:", temp_tensor.dtype)
```

```
Adjustment shape: torch.Size([2, 3])
Adjustment type: torch.int64
Temperatures shape: torch.Size([2, 3])
Temperatures type: torch.int64
```

Instructions 2

- Add the `temperatures` and `adjustment` tensors.

```
adjustment = torch.tensor([[2, 2, 2], [2, 2, 2]])

# Add the temperatures and adjustment tensors
corrected_temperatures = temp_tensor + adjustment
print("Corrected temperatures:", corrected_temperatures)
```

```
Corrected temperatures: tensor([[74, 77, 80],
                               [72, 75, 78]])
```

Exercise 3

Linear layer network

Neural networks often contain many layers, but most of them are linear layers. Understanding a single linear layer helps you grasp how they work before adding complexity.

Apply a `linear layer` to an input tensor and observe the output.

Instructions

- Create a Linear layer that takes 3 features as input and returns 2 outputs.
- Pass `input_tensor` through the linear layer.

```
#Import neural network from torch library
#import torch library when use in the exercise
import torch.nn as nn

input_tensor = torch.tensor([[0.3471, 0.4547, -0.2356]])

# Create a Linear layer
linear_layer = nn.Linear(
    in_features=3,
    out_features=2
)

# Pass input_tensor through the linear layer
output = linear_layer(input_tensor)

print(output)

tensor([[ -0.4160,  0.6729]], grad_fn=<AddmmBackward0>)
```

Exercise 4

Exercise 5

Your first neural network

It's time for you to implement a small neural network containing two linear layers in sequence.

Instructions

- Add a container for stacking layers in sequence.

```
# Import libraries when use in the exercise

input_tensor = torch.Tensor([[2, 3, 6, 7, 9, 3, 2, 1]])

# Create a container for stacking linear layers
model = nn.Sequential(nn.Linear(8, 4),
    nn.Linear(4, 1)
)

output = model(input_tensor)
print(output)

tensor([[0.9745]], grad_fn=<AddmmBackward0>)
```

Exercise 5

Exercise 6

Counting the number of parameters

Deep learning models are famous for having a lot of parameters. With more parameters comes more computational complexity and longer training times, and a deep learning practitioner must know how many parameters their model has.

In this exercise, you'll first calculate the number of parameters manually. Then, you'll verify your result using the `.numel()` method.

Instructions 1/2

Question Manually calculate the number of parameters of the model below. How many does it have? Use the console as a calculator.

```
model = nn.Sequential(nn.Linear(9, 4),  
                      nn.Linear(4, 2),  
                      nn.Linear(2, 1))
```

[Copy](#)

Possible answers

1. 53 (correct)

2. 16

3. 29

Instructions 2/2

- Use `.numel()` to confirm your manual calculation by iterating through the model's parameters to updating the total variable.

```
import torch.nn as nn  
  
model = nn.Sequential(nn.Linear(9, 4),  
                      nn.Linear(4, 2),  
                      nn.Linear(2, 1))  
  
total = 0  
  
# Calculate the number of parameters in the model  
for p in model.parameters():  
    total += p.numel()  
  
print(f"The number of parameters in the model is {total}")
```

The number of parameters in the model is 53